

# Servant


## Haskell and Cryptocurrencies

---

Dr. Andres Löh, Well-Typed LLP

Dr. Lars Brünjes, IOHK

2018-03-05



INPUT | OUTPUT

# Goals

- A case study for type-level programming.
- An API description language.
- Using **servant** to define clients and servers.

# Describing an API

---

# A simple web API

GET	/entry	get all entries
POST	/entry	post a new entry
GET	/entry/<n>	get entry of given number
DELETE	/entry/<n>	remove entry of given number

## Describing the API as a Haskell type

```
type KeyVal =  
    "entry" :> Get '[JSON] [Entry]  
:<|> "entry"  
    :> ReqBody '[JSON] Text  
    :> PostCreated '[JSON] Entry  
:<|> "entry" :> Capture "n" Int  
    :> Get '[JSON] Entry  
:<|> "entry" :> Capture "n" Int  
    :> DeleteNoContent '[JSON] NoContent
```

## Describing the API as a Haskell type

```
type KeyVal =  
    "entry" :> Get '[JSON] [Entry]  
  :<|> "entry"  
      :> ReqBody '[JSON] Text  
      :> PostCreated '[JSON] Entry  
  :<|> "entry" :> Capture "n" Int  
      :> Get '[JSON] Entry  
  :<|> "entry" :> Capture "n" Int  
      :> DeleteNoContent '[JSON] NoContent
```

Let's first look at how this is even kind-correct.

## API combinators

```
data a :<|> b
infixr 8 :<|>

data path :> a
infixr 9 :>

data Capture (sym :: Symbol) a
data ReqBody (contentType :: [*]) a

type Get = Verb GET 200
type PostCreated = Verb POST 201
type DeleteNoContent = Verb DELETE 204

data Verb
  method (status :: Nat) (contentType :: [*]) a
```

Implementations are irrelevant.

# Meaning

- $a :<|> b$  is the union of the APIs  $a$  and  $b$ ,



# Meaning

- `a :<|> b` is the union of the APIs `a` and `b`,
- `"foo" :> a` describes an API reachable under path `"foo"`,

# Meaning

- `a :<|> b` is the union of the APIs `a` and `b`,
- `"foo" :> a` describes an API reachable under path `"foo"`,
- `Get '[JSON] Entry` describes an endpoint returning an `Entry` with content type `JSON`,

# Meaning

- `a :<|> b` is the union of the APIs `a` and `b`,
- `"foo" :> a` describes an API reachable under path `"foo"`,
- `Get '[JSON] Entry` describes an endpoint returning an `Entry` with content type `JSON`,
- `Capture "n" Int :> a` describes an API reachable under a captured integer path,

# Meaning

- `a :<|> b` is the union of the APIs `a` and `b`,
- `"foo" :> a` describes an API reachable under path `"foo"`,
- `Get '[JSON] Entry` describes an endpoint returning an `Entry` with content type `JSON`,
- `Capture "n" Int :> a` describes an API reachable under a captured integer path,
- `ReqBody '[JSON] Text :> a` describes an API reachable only if an appropriate request body is present.

## More API combinators

```
data JSON
```

```
data PlainText
```

```
data StdMethod =
```

```
    GET | POST | HEAD | PUT | DELETE  
    | TRACE | CONNECT | OPTIONS | PATCH
```

```
data NoContent = NoContent
```

## Overloaded strings and numbers

```
GHCi> import GHC.TypeLits
GHCi> :k "foo"
"foo" :: Symbol
GHCi> :k 42
42 :: Nat
```

# Entries

Unlike all the others, `Entry` is a domain-specific type for our API and to be defined in our program:

```
data Entry =  
  Entry  
    { entryId    :: Int  
    , entryText  :: Text  
    }
```

# How to model a web server

---



# Request-response

At a high-level, a web server receives **requests** and sends **responses**.

# Request-response

At a high-level, a web server receives **requests** and sends **responses**.

The `wai` package (short for **web application interface**) is built on top of this simple model:

```
type Application =  
    Request  
    -> (Response -> IO ResponseReceived)  
    -> IO ResponseReceived
```

(This is in *continuation passing style*.)

# Running an application

From `Network.Wai.Handler.Warp`:

```
run :: Port -> Application -> IO ()
```

Starts a webserver and runs the given application on the given port.

## A trivial webserver

We can build an application that ignores the request and always answers with a `200` response and the text `"Hello"`:

```
alwaysHello :: Application
alwaysHello req respond =
    respond (responseLBS status200 [] "Hello\n")
```

```
test :: IO ()
test =
    run 8888 alwaysHello
```

```
$ curl -X GET "http://localhost:8888/foo/bar"  
Hello  
$ curl -X DELETE "http://localhost:8888/foo/baz"  
Hello
```

# Routing an dispatching

What a “real” application has to do:

- parse the incoming requests,
- depending on the path, the request headers, and the request body, make decisions on how to handle the request,
- dispatch the request to an appropriate handler,
- take the response from the handler and forward it to the request sender, or produce an error response instead.

# Routing an dispatching

What a “real” application has to do:

- parse the incoming requests,
- depending on the path, the request headers, and the request body, make decisions on how to handle the request,
- dispatch the request to an appropriate handler,
- take the response from the handler and forward it to the request sender, or produce an error response instead.

These tasks can be automated using **servant** once the API is known.

# Computing an application

---



## The interface

```
serve ::  
  HasServer api '[]  
  => Proxy api -> Server api -> Application
```

# The interface

```
serve ::  
  HasServer api '[]  
  => Proxy api -> Server api -> Application
```

- `api` is an API type such as we have seen,

# The interface

```
serve ::  
  HasServer api '[]  
  => Proxy api -> Server api -> Application
```

- `api` is an API type such as we have seen,
- `HasServer` is a type class defined by `servant`,

# The interface

```
serve ::  
  HasServer api '[]  
  => Proxy api -> Server api -> Application
```

- `api` is an API type such as we have seen,
- `HasServer` is a type class defined by `servant`,
- `Server` is a type family defined by `servant` (non-injective, therefore the `Proxy`).

## Server and HasServer

Actually, `Server` is a wrapper around an associated type:

```
class HasServer api (context :: [*]) where
  type SeverT api (m :: * -> *) :: *
  route ::
    Proxy api
    -> Context context
    -> Delayed env (Server api)
    -> Router env
```

```
type Server api = SeverT api Handler
newtype Handler a = Handler
  {runHandler' :: ExceptT ServantErr IO a}
```

# Simplify

We simplify a bit in order to better keep track of the actual ideas:

- We move `SeverT` out of the `HasServer` class.
- We ignore “contexts” for now.
- We ignore `Delayed` and `Router` and adopt a simpler model instead.

## Simplified scenario

```
type family ServerT api (m :: * -> *) :: *
class HasServer api where
  route ::
    Proxy api
    -> Server api
    -> Request
    -> Handler Response
```

## Simplified scenario

```
type family ServerT api (m :: * -> *) :: *  
class HasServer api where  
  route ::  
    Proxy api  
    -> Server api  
    -> Request  
    -> Handler Response
```

Intuitively, `ServerT` tells us the types of the handlers we need for a given API, and `route` turns the given handlers into a request-response system.



## Implementing `ServerT` – verbs

```
type instance ServerT  
  (Verb method status ctypes a) m =  
  m a
```

Recall that `m` is instantiated to `Handler`, and `Handler a` is `ExceptT ServantErr IO a`.

So at an endpoint, we simply provide an `IO` action with the capability to throw errors and the obligation to return a value of result type `a`.

## Implementing `ServerT` – paths

```
type instance ServerT  
  ((path :: Symbol) => api) m =  
  ServerT api m
```

A path does not affect the type of handler at all.

## Implementing `ServerT` – captures

```
type instance ServerT  
  (Capture capture a := api) m =  
  a -> ServerT api m
```

A capture of type `a` is provided to the handler as an *additional input* of type `a`!

## Implementing `ServerT` – request body

```
type instance ServerT  
  (ReqBody ctypes a :=> api) m =  
  a -> ServerT api m
```

The same happens for the request body which can also be accessed in the handler.

## Implementing `ServerT` – choice

```
type instance ServerT  
  (api1 :<|> api2) m =  
  ServerT api1 m :<|> ServerT api2 m
```

```
data a :<|> b = a :<|> b  
infixr 8 :<|>
```

A handler for two combined APIs is the combination (i.e., a pair) of the handlers for the two individual APIs.

We use a Servant-specific pair constructor (just for optical reasons).

## Applying the type family

We can now (e.g. in GHCi) compute `Server KeyVal` and obtain:

```
Server KeyVal
~ (      Handler [Entry]                -- get all
  :<|> (Text -> Handler Entry)         -- post
  :<|> (Int  -> Handler Entry)         -- get single
  :<|> (Int  -> Handler NoContent)     -- delete
)
```

## Implementing the handlers

---

## Maintaining server state

We keep the server state in an STM `TVar`:

```
type Store = TVar (Map Int Text)
```



## Maintaining server state

We keep the server state in an STM `TVar`:

```
type Store = TVar (Map Int Text)
```

This gives us thread safety – because when a `wai` application is run, every request will be handled in its own Haskell thread, and be concurrent with possibly many others.

## Maintaining server state

We keep the server state in an STM `TVar`:

```
type Store = TVar (Map Int Text)
```

This gives us thread safety – because when a `wai` application is run, every request will be handled in its own Haskell thread, and be concurrent with possibly many others.

In a proper system we would probably use a database to make the state persistent.

## Getting all entries

```
getAllEntries :: Store -> Handler [Entry]
getAllEntries store = do
  entries <- liftIO $ readTVarIO store
  return $ map (uncurry Entry) (toList entries)
```

## Posting a new entry

```
postEntry :: Store -> Text -> Handler Entry
postEntry store txt =
  liftIO $ atomically $ do
    entries <- readTVar store
    let key = if M.null entries then 1
              else fst (findMax entries) + 1
    writeTVar store (insert key txt entries)
    return (Entry key txt)
```

## Getting a single entry

```
getEntry :: Store -> Int -> Handler Entry
getEntry store n = do
  entries <- liftIO $ readTVarIO store
  case lookup n entries of
    Just txt -> return (Entry n txt)
    Nothing   -> throwError err404
```

## Deleting an entry

```
deleteEntry :: Store -> Int -> Handler NoContent
deleteEntry store n = do
  liftIO $ atomically $
    modifyTVar store (delete n)
  return NoContent
```

# Everything together

```
keyVal :: Store -> Server KeyVal
keyVal store =
    getAllEntries store
  :<|> postEntry store
  :<|> getEntry store
  :<|> deleteEntry store
```

This is (nearly) all we have to do in order to be able to run the server.

## Completing the implementation

---



## The `HasServer` class

Let's discuss how we can implement our (simplified) `HasServer` class:

```
class HasServer api where
  route ::
    Proxy api
    -> Server api
    -> Request
    -> Handler Response
```

```
data Verb  
  method (status :: Nat) (contentType :: [*]) a
```

## Implementing `HasServer` – verbs

```
data Verb
```

```
  method (status :: Nat) (contentType :: [*]) a
```

The `method` indicates which requests match. We need a way to turn the `method` type into a value.

## Implementing `HasServer` – verbs

```
data Verb
```

```
  method (status :: Nat) (contentType :: [*]) a
```

The `method` indicates which requests match. We need a way to turn the `method` type into a value.

The `status` indicates which response code to return. We need a way to turn the `status` type into a value.

## Implementing `HasServer` – verbs

```
data Verb
```

```
  method (status :: Nat) (contentType :: [*]) a
```

The `method` indicates which requests match. We need a way to turn the `method` type into a value.

The `status` indicates which response code to return. We need a way to turn the `status` type into a value.

The `contentType` indicate which formats we can convert our output to. We need some way to translate `a` into these formats.

## Reflecting methods

```
class ReflectMethod a where  
  reflectMethod :: Proxy a -> Method
```

## Reflecting methods

```
class ReflectMethod a where  
  reflectMethod :: Proxy a -> Method
```

```
instance ReflectMethod GET where  
  reflectMethod _ = methodGet
```

```
instance ReflectMethod POST where  
  reflectMethod _ = methodPost
```

```
instance ReflectMethod DELETE where  
  reflectMethod _ = methodDelete
```

## Reflecting numbers (and symbols)

This is provided by GHC itself:

```
natVal      :: KnownNat n => Proxy n -> Integer
symbolVal   :: KnownSymbol n => Proxy n -> String
```

The constraints `KnownNat` and `KnownSymbol` are satisfied for all literal numbers and strings.



## Reflecting numbers (and symbols)

This is provided by GHC itself:

```
natVal      :: KnownNat n => Proxy n -> Integer
symbolVal   :: KnownSymbol n => Proxy n -> String
```

The constraints `KnownNat` and `KnownSymbol` are satisfied for all literal numbers and strings.

```
GHCi> natVal (Proxy @ 42)
42
GHCi> symbolVal (Proxy @ "foo")
"foo"
```

# Content types

We again present a slightly simplified view:

```
class AllCTRender (ctypes :: [*]) a where
  handleAcceptH ::
    Proxy ctypes -> AcceptHeader
    -> Maybe (ContentType, a -> ByteString)
```

# Content types

We again present a slightly simplified view:

```
class AllCTRender (ctypes :: [*]) a where
  handleAcceptH ::
    Proxy ctypes -> AcceptHeader
    -> Maybe (ContentType, a -> ByteString)
```

The `handleAcceptH` method takes an accept header (from a request) and a value and decides which content type and representation to use (or fails).

## Base case

```
instance AllCTRender '[] a where  
  handleAcceptH _ _ = Nothing
```

If there are no accepted content types, then we can never succeed.

## Case for **JSON**

```
instance (ToJSON a, AllCTRender cs a)
  => AllCTRender (JSON ': cs) a where
  handleAcceptH _ (AcceptHeader accept) =
    let
      ct = "application/json"
    in
      (ct, encode) <$ matchAccept [ct, "*/*"] accept
      <|> handleAcceptH (Proxy @ cs)
        (AcceptHeader accept)
```

We check if **"application/json"** is among the accepted content types, and either perform the rendering (requiring a **ToJSON a** instance) or fall back on other content types.

- Check that the request method matches.
- Check that the request path is empty.
- Check that the accept header is compatible.
- If all this applies, run the handler and encode the result according to the selected content type.
- In any other case, fail.

## Implementing `HasServer` – verbs

```
instance
  ( AllCTRender ctypes a
  , ReflectMethod method, KnownNat status)
=> HasServer (Verb method status ctypes a) where
route _ h req
  | requestMethod req == reflectMethod (Proxy @ method)
  && null (pathInfo req) =
    case handleAcceptH (Proxy @ ctypes)
      (acceptHeader req) of
      Just (ct, enc) -> do
        a <- h
        return $ responseLBS
          (Status
            (fromIntegral (natVal (Proxy @ status)))) ""
          [("Content-Type", ct)]
          (enc a)
        Nothing -> throwError err404 -- drastic simplification
  | otherwise = throwError err404 -- drastic simplification
```

## Implementing `HasServer` – paths

```
instance
  (KnownSymbol path, HasServer api)
  => HasServer (path :> api) where
  route _ h req =
    case pathInfo req of
      (x : xs)
        | x == fromString (symbolVal (Proxy @ path)) ->
          route (Proxy @ api) h (req {pathInfo = xs})
      _ -> throwError err404
```

Comparatively simple – we just check if the path prefix matches.



## Implementing `HasServer` – captures

```
data Capture (sym :: Symbol) a
```

For captures, we have to be able to parse a path component according to the type of the capture `a`.

The `Symbol` is not being used here – it is contained in the API purely for documentation purposes.

# Parsing captures

Once again, we take a simplified view:

```
class FromHttpApiData a where  
  parseUrlPiece :: Text -> Maybe a
```

```
instance FromHttpApiData Int where  
  parseUrlPiece = readMaybe . unpack
```

## Implementing `HasServer` – captures

```
instance
  (FromHttpApiData a, HasServer api)
  => HasServer (Capture capture a :=> api) where
  route _ h req =
    case pathInfo req of
      (x : xs) ->
        case parseUrlPiece x of
          Just a ->
            route (Proxy @ api) (h a) (req {pathInfo = xs})
          Nothing -> throwError err404
      _ -> throwError err404
```

Note how we pass the parsed path fragment to the handler when continuing.

For the request body, we have to do content type negotiation once more – but in the other direction:

The content type of the body must match one of the content types we are able to parse.

# Content types

```
class AllCTUnrender (ctypes :: [*]) a where
  canHandleCTypeH ::
    Proxy ctypes
    -> ContentType
    -> Maybe (ByteString -> Maybe a)
```

The outer `Maybe` signals whether we can handle the given content type – the inner `Maybe` signals parse failure.

## Base case

```
instance AllCTUnrender '[] a where  
  canHandleCTypeH _ _ = Nothing
```

Without any accepted content types we unsurprisingly fail.

```
instance (FromJSON a, AllCTUnrender cs a)
  => AllCTUnrender (JSON ': cs) a where
  canHandleCTypeH _ ct
    | ct == "application/json" = Just decode
    | otherwise                =
      canHandleCTypeH (Proxy @ cs) ct
```

## Implementing `HasServer` – request body

```
instance
  (AllCTUnrender ctypes a, HasServer api)
  => HasServer (ReqBody ctypes a :> api) where
  route _ h req =
    case canHandleCTypeH (Proxy @ ctypes)
      (contentTypeHeader req) of
      Just dec -> do
        b <- liftIO $ lazyRequestBody req
        case dec b of
          Just a  -> route (Proxy @ api) (h a) req
          Nothing -> throwError err404
      Nothing -> throwError err404
```



The final case we have to implement is choice:

- In our simple setting, we test if the first alternative fails with 404, and only if it does, try the second one.
- Full `servant` has a much more sophisticated notion of error codes and priorities, to make sure that the “correct” response code is sent in most situations.

## Implementing `HasServer` – choice

```
instance
  (HasServer api1, HasServer api2)
  => HasServer (api1 :<|> api2) where
route _ (h1 :<|> h2) req =
  catchError
    (route (Proxy @ api1) h1 req)
    (\ err ->
      if errHTTPCode err == 404
      then route (Proxy @ api2) h2 req
      else throwError err)
```

## What is left to do?

We still need to implement `serve` in terms of `route`:

```
serve ::  
  HasServer api  
    => Proxy api -> Server api -> Application  
serve p h req respond = do  
  result <-  
    runExceptT (runHandler' (route p h req))  
  case result of  
    Left err  -> respond (responseServantErr err)  
    Right rsp -> respond rsp
```

## Entry and JSON

We also still need to provide `FromJSON` and `ToJSON` instances for `Entry` – a trivial job thanks to generic programming:

```
deriving instance Generic Entry
instance FromJSON Entry
instance ToJSON Entry
```

## Putting everything together

```
main :: IO ()
main = do
  store <- newTVarIO empty
  run 8000
    (serve (Proxy @ KeyVal) (keyVal store))
```

This allocates a `TVar` to hold the store, then computes the server from the API and the handlers we defined before (`keyVal`), and finally runs it.

# Testing

```
$ curl -X GET "http://localhost:8000/entry"
[]
$ curl -H "Content-Type: application/json" -d "\"foo\""
-X POST "http://localhost:8000/entry"
{"entryId":1,"entryText":"foo"}
$ curl -H "Content-Type: application/json" -d "\"bar\""
-X POST "http://localhost:8000/entry"
{"entryId":2,"entryText":"bar"}
$ curl -X GET "http://localhost:8000/entry"
[{"entryId":1,"entryText":"foo"}, {"entryId":2,"entryText":"bar"}]
$ curl -X GET "http://localhost:8000/entry/1"
{"entryId":1,"entryText":"foo"}
$ curl -X DELETE "http://localhost:8000/entry/1"
$ curl -X GET "http://localhost:8000/entry"
[{"entryId":2,"entryText":"bar"}]
```

# Recap

We've spent a lot of time on basically re-implementing **servant**, so let's briefly look at what we actually had to do.

We had to:

- Define the API.
- Define the domain datatypes and instances (`Entry`, `FromJSON`, `ToJSON`).
- Define the handlers.
- Compose everything (`main`).

```
type KeyVal =  
  "entry" :=> Get '[JSON] [Entry]  
:<|> "entry"  
  :=> ReqBody '[JSON] Text  
  :=> PostCreated '[JSON] Entry  
:<|> "entry" :=> Capture "n" Int  
  :=> Get '[JSON] Entry  
:<|> "entry" :=> Capture "n" Int  
  :=> DeleteNoContent '[JSON] NoContent
```



## Domain datatypes and instances

```
data Entry =  
  Entry  
    { entryId    :: Int  
    , entryText  :: Text  
    }  
  deriving (Generic, Show)  
instance FromJSON Entry  
instance ToJSON Entry
```

# Handlers

```
getAllEntries :: Store -> Handler [Entry]
getAllEntries store = do
  entries <- liftIO $ readTVarIO store
  return $ map (uncurry Entry) (toList entries)

postEntry :: Store -> Text -> Handler Entry
postEntry store txt =
  liftIO $ atomically $ do
    entries <- readTVar store
    let key = if M.null entries then 1
              else fst (findMax entries) + 1
    writeTVar store (insert key txt entries)
    return (Entry key txt)

getEntry :: Store -> Int -> Handler Entry
getEntry store n = do
  entries <- liftIO $ readTVarIO store
  case lookup n entries of
    Just txt -> return (Entry n txt)
    Nothing   -> throwError err404

deleteEntry :: Store -> Int -> Handler NoContent
deleteEntry store n = do
  liftIO $ atomically $
    modifyTVar store (delete n)
  return NoContent
```

## Combined handler

```
keyVal :: Store -> Server KeyVal
keyVal store =
    getAllEntries store
  :<|> postEntry store
  :<|> getEntry store
  :<|> deleteEntry store
```

## Composing everything

```
main :: IO ()  
main = do  
  store <- newTVarIO empty  
  run 8000  
    (serve (Proxy @ KeyVal) (keyVal store))
```

# What do we get?

A lot of code “for free”:

- All the routing, parsing, serialization, dispatching is taken care of automatically, via the API specification.
- The handlers are just ordinary functions that do not need to know much about running in a web setting.

Type safety:

- The handlers work with dedicated Haskell types, all the conversions to less typed settings (JSON, strings) happen behind the scenes.
- If we change the API type, or use an undeclared input, or produce the wrong output, this is all a type error.

Other “interpretations”

---

# Servant is a domain-specific language

What **servant** really provides is a type-level domain-specific language for describing APIs:

- The API is the “program”.
- The server is an “interpretation” of the program.

# Servant is a domain-specific language

What **servant** really provides is a type-level domain-specific language for describing APIs:

- The API is the “program”.
- The server is an “interpretation” of the program.

But **servant** offers more interpretations besides the server interpretation.



## Deriving a client

---

# Computing client functions

```
client ::  
  HasClient api  
  => Proxy api -> Client api
```

Very similar – `HasClient` is a class, and `Client` is an associated type.

## Instantiating `Client`

```
Client KeyVal
~ (      ClientM [Entry]           -- get all
   :<|> (Text -> ClientM Entry)    -- post
   :<|> (Int  -> ClientM Entry)    -- get single
   :<|> (Int  -> ClientM NoContent) -- delete
)
```

Nearly the same as `Server`, only that this time, we *get* these functions, rather than having to provide them.

## Obtaining client functions

```
getAll      :: ClientM [Entry]
postNew     :: Text -> ClientM Entry
getSingle   :: Int -> ClientM Entry
delSingle   :: Int -> ClientM NoContent

getAll :<|> postNew :<|> getSingle :<|> delSingle
      = client (Proxy @ KeyVal)
```

## Running client functions

```
runClientM ::  
  ClientM a -> ClientEnv  
  -> IO (Either ServantError a)
```

## Running client functions

```
runClientM ::  
  ClientM a -> ClientEnv  
  -> IO (Either ServantError a)
```

```
data ClientEnv =  
  ClientEnv  
    { manager :: Manager  
    , baseUrl :: BaseUrl  
    }
```

## Testing in GHCi

Of course, this requires we are running the server.

We can create a manager using the `http-client` package.

```
GHCi> m <- newManager defaultManagerSettings
GHCi> base = BaseUrl Http "localhost" 8000 ""
GHCi> env = ClientEnv m base
GHCi> runClientM (postNew "foo") env
Right (Entry {entryId = 1, entryText = "foo"})
GHCi> runClientM getAll env
Right [Entry {entryId = 1, entryText = "foo"}]
```

If you derive both a server and a client from a common API, you can be certain that they match.



## Servant for clients

If you derive both a server and a client from a common API, you can be certain that they match.

On the other hand, it can also be useful to employ **servant** to just derive a client for a web API you want to bind to, in order to quickly get the right functions in a strongly typed way.

## Yet more interpretations

There's more:

- Generate clients in various other / external languages (e.g. Javascript).
- Generate documentation in various formats (e.g. Swagger).
- Generate type-safe links within an API.
- Generate mock servers and mock clients.
- ...

Servant is carefully designed to be extensible in two dimensions:

- You can add new interpretations.
- You can add new combinators to the API language.

Both kinds of extension usually do not require access or modifications of the core library.